

```
In [6]: import numpy as np
from sklearn import linear_model, datasets, cross_validation
iris = datasets.load_iris()
X = iris.data
y = iris.target
X_train, X_test, y_train, y_test = cross_validation.train_test_split(X, y, test_size=0.3)
clf = linear_model.LogisticRegression()
clf.fit(X_train, y_train)
accuracy = clf.score(X_test, y_test)
print(accuracy)

0.911111111111
```

```
In [12]: len(X_train)
```

```
Out[12]: 105
```

```
In [13]: len(y_train)
```

```
Out[13]: 105
```

```
In [14]: len(X_test)
```

```
Out[14]: 45
```

```
In [15]: len(y_test)
```

```
Out[15]: 45
```

```
In [ ]:
```